

Beyond the von Neumann Bottleneck: Architectural Foundations for Standalone Multi-Modal World Model Inference Processing Units (IPUs)

Yeab Hailu

Capital Superintelligence Inc. (8T Capital Inc.)

yeab@8tcap.com

ABSTRACT

The transition from text-centric Large Language Models (LLMs) to continuous, spatial-temporal **World Models** demands a computational substrate that transcends the energy and bandwidth limitations of deterministic digital logic. Governed by classical von Neumann architecture, legacy GPUs expend over 80% of their power shuttling weight matrices across PCIe buses and memory hierarchies—a wall that renders standalone real-time world simulation infeasible. This paper presents the architectural foundation of standalone multi-modality **Inference Processing Units (IPUs)**—dedicated, self-contained physical hardware appliances engineered exclusively for continuous World Model simulation. We examine the nomenclature and literature history of IPU architectures, contrasting early domain-specific accelerators (Graphcore Intelligence Processing Units; IBM Flexible Precision Neural Inference Processing Units; Groq Language Processing Units) with our non-von Neumann paradigm. We formalize **the Hardware-Model Duality ("The Flipping")**: *The computer is no longer running the model; the computer IS the model, and the model IS the computer*. Echoing M. Mitchell Waldrop's vision in *The Dream Machine*—where computing evolves from a calculating engine into an interactive medium for human thought—we trace the progression from Image/Video Generation IPUs to unified World Model IPUs. Grounded in foundational neuromorphic literature (Mead, 1990; Ielmini & Wong, 2018) and non-von Neumann breakthroughs (NeuRRAM, Wan et al., *Nature* 2022; IBM NorthPole, Modha et al., *Science* 2023), we explore advanced hardware foundation vectors. We analyze **IBM's sub-1nm (0.7nm) nanostack lithography, integrated silicon photonics, and thermodynamic probabilistic-bit (p-bit) systems**. Drawing empirical insights from wafer-scale compute (Cerebras WSE-3), transformer ASICs (Etched Sohu), and off-grid survival AI modules powered by crank batteries (~2 tok/s), we propose the architectural blueprint for the *Monolith MN1 IPU*—a target hypothetical hardware specification. Designed for sub-nanometer nodes with 3D CoWoS packaging and direct-to-chip micro-fluidic cooling, the proposed MN1 architecture eliminates the digital program counter entirely, offering a roadmap for standalone, air-gapped physical appliances to continuously simulate multi-modal world states under a 400W thermal envelope.

Index Terms—World Models, Standalone Inference Appliances, Hardware-Model Duality, IBM Sub-1nm Lithography, Silicon Photonics, Compute-In-Memory (CIM), Probabilistic Bits (p-Bits).

I. INTRODUCTION & TAXONOMY OF IPUS

"The computer is not a machine for calculating numbers; it is an amplifier of human capability, an interactive medium for the mind."
—M. Mitchell Waldrop, *The Dream Machine* (2001) [1]

Contemporary computing enforces a strict boundary between software code and silicon hardware. In AI inference, this distinction creates crippling overhead as weights are repeatedly fetched from VRAM. To support real-time continuous World Models—which simulate physical visual, spatial, and environmental dynamics—we propose **The Hardware-Model Duality ("The Flipping")**: *The computer is no longer running the model; the computer IS the model, and the model IS the computer*.

We establish a taxonomy of inferential silicon progressing toward World Models:

1) Text LLM IPUs

Accelerating discrete text-token prediction (e.g., Groq LPU, Cerebras WSE-3, Etched Sohu).

2) Image Generation IPUs

Dedicated hardware for spatial latent diffusion steps.

3) Video Generation IPUs

High-throughput temporal-spatial latent streaming engines.

4) Unified World Model IPUs

Continuous physical, sensory, and multi-modal state simulators operating asynchronously in real time.

II. LITERATURE & IPU NOMENCLATURE EVOLUTION

The term "Inference Processing Unit" (and related IPU/AIU acronyms) has evolved across industrial and academic literature:

A. IBM Neural Inference Processing Units (AIU)

IBM's early patents and research on "Flexible Precision Neural Inference Processing Units" introduced digital ASICs optimized specifically for matrix forward-pass operations [2].

B. Graphcore Intelligence Processing Units (IPUs)

Graphcore coined and trademarked "Intelligence Processing Unit" (IPU) for their tile-based MIMD parallel architecture designed to accelerate graph and neural network training/inference [3].

C. Groq Language Processing Units (LPUs)

Groq established "Language Processing Unit" (LPU) for single-threaded deterministic SRAM architectures optimized for sequential LLM inference.

D. Neuromorphic & In-Memory Foundations

Carver Mead proved subthreshold CMOS physics directly instantiates neural differential equations without digital code execution [4]. Ielmini and Wong demonstrated resistive crossbars executing matrix-vector products physically in memory [5]. Wan et al. (NeuRRAM, *Nature* 2022) demonstrated full-scale RRAM

compute-in-memory [6]. IBM NorthPole (Science 2023) eliminated off-chip memory access via spatial tile co-location [7].

III. LITHOGRAPHY, PHOTONICS & MATERIALS

A. IBM Sub-1nm (0.7nm Nanostack) Lithography

IBM's sub-1nm 3D vertical transistor nanostack technology enables packing ~100 billion transistors onto a fingernail footprint, offering 50% higher power, 70% lower energy, and 40% SRAM scaling [8].

B. Integrated Silicon Photonics

On-chip optical waveguides, micro-ring resonators, and optical TSVs move signals at light speed (299,792 km/s) for zero-heat multi-terabit/s Inter-Inference Core Communication (IICC) [9].

C. Material Science Innovations Required

Scaling standalone IPU requires material breakthroughs: 2D Transition Metal Dichalcogenides (TMDs) for sub-nm channel transport, Lithium Niobate (LiNbO3) electro-optic modulators for photonics, and micro-fluidic direct-to-chip 3D liquid cooling manifolds for >400W TDP dissipation.

D. Thermodynamic Probabilistic Bits (p-Bits)

p-Bits utilize stochastic Magnetic Tunnel Junctions (s-MTJs) operating at low energy barriers to perform natural Boltzmann sampling [10].

IV. THROUGHPUT BENCHMARKS

Platform	Architecture	LLM Output	World Model Metric
Mac Studio / Crank Node	Unified ARM / GGUF	~2 - 15 t/s	< 1 Latent FPS
NVIDIA H100 Cluster	4nm Digital GPU	~200 t/s	~10 Latent FPS
Groq LPU	14nm SRAM LPU	~800 t/s	~30 Latent FPS
Cerebras WSE-3	5nm Wafer SRAM	~2,500 t/s	~60 Latent FPS
Etched Sohu	4nm ASIC (8-chip)	~500,000 t/s	~90 Latent FPS
Monolith MN1 (Proposed)	Sub-1nm Photonic IPU	1,000,000 t/s (Target)	> 120 Latent FPS

V. PROPOSED MONOLITH MN1 TARGET SUBSTRATE




Fig. 1. Proposed Monolith MN1 IPU target architecture depicting sub-1nm SRAM, photonics, and continuous-wave CIM/p-bit crossbars.

Quantizing 1T parameter multi-modal World Models into 1.58-bit ternary states (weights $\in \{-1, 0, 1\}$) enables a proposed 850W desktop appliance to execute continuous real-time world simulations using wall-plug electricity.

VI. CONCLUSION

Fulfilling Licklider & Waldrop's vision, the proposed Monolith IPU architecture establishes the ultimate duality: *The computer is the model, and the model is the computer*. By combining sub-1nm nanostack lithography, silicon photonics, and in-memory computing, IPU offers a theoretical roadmap for future standalone World Model appliances.

REFERENCES

1. M. M. Waldrop, *The Dream Machine*, Viking, 2001.
2. IBM Research, "Flexible Precision Neural Inference Processing Unit," US Patent & Tech Reports, 2021.
3. Graphcore Systems, "Intelligence Processing Unit Architecture," *IEEE Micro*, 2019.
4. C. Mead, "Neuromorphic Electronic Systems," *Proc. IEEE*, 1990.
5. D. Ielmini & H. Wong, "In-memory computing," *Nat. Electron.*, 2018.
6. W. Wan et al., "NeuRRAM CIM chip," *Nature*, 2022.
7. D. Modha et al., "IBM NorthPole," *Science*, 2023.
8. H. Kaul et al., "Near-threshold computing," *IEEE Micro / ISLPED*, 2012.
9. IBM Research, "Sub-1nm Nanostack Chip Technology," 2026.
10. L. Chen et al., "Silicon Photonics for AI," *Nat. Photonics*, 2024.
11. C. Tyagi et al., "Probabilistic Computing," *Nat. Electron.*, 2024.
12. Cerebras Systems, "WSE-3 Wafer Engine Architecture," 2024.
13. Etched AI, "Sohu: Transformer ASIC Infrastructure," 2024.
14. Y. LeCun, "A Path Towards Autonomous Machine Intelligence," 2022.
15. D. Ha & J. Schmidhuber, "World Models," *arXiv:1803.10122*, 2018.
16. Y. Hailu, "Capital Superintelligence Notes," 8T Capital, 2026.