

Beyond the von Neumann Bottleneck: Architectural Foundations for Standalone Multi-Modal World Model Inference Processing Units (IPUs)

Yeab Hailu

Capital Superintelligence Inc. (8T Capital Inc.)

yeab@8tcap.com

ABSTRACT

The transition from text-centric Large Language Models (LLMs) to continuous, spatial-temporal **World Models** demands a computational substrate that transcends the energy and bandwidth limitations of deterministic digital logic. Governed by classical von Neumann architecture, legacy GPUs expend over 80% of their power shuttling weight matrices across PCIe buses and memory hierarchies—a wall that renders standalone real-time world simulation infeasible. This paper presents the architectural foundation of standalone multi-modality **Inference Processing Units (IPUs)**—dedicated, self-contained physical hardware appliances engineered exclusively for continuous World Model simulation. While LLM IPUs represent an essential early sub-category of inferential silicon, true intelligence requires scaling beyond discrete text token generation toward high-frame-rate spatial, temporal, and physical world simulation. We formalize **the Hardware-Model Duality ("The Flipping")**: *The computer is no longer running the model; the computer IS the model, and the model IS the computer.* Echoing M. Mitchell Waldrop's vision in *The Dream Machine*—where computing evolves from a calculating engine into an interactive medium for human thought—we trace the progression from Image/Video Generation IPUs to unified World Model IPUs. Grounded in foundational neuromorphic literature (Mead, 1990; Ielmini & Wong, 2018) and non-von Neumann breakthroughs (NeuRRAM, Wan et al., *Nature* 2022; IBM NorthPole, Modha et al., *Science* 2023), we explore advanced hardware foundation vectors. We analyze **IBM's sub-1nm (0.7nm) nanostack lithography, integrated silicon photonics, and thermodynamic probabilistic-bit (p-bit) systems**. Drawing empirical insights from wafer-scale compute (Cerebras WSE-3), transformer ASICs (Etched Sohu), and off-grid survival AI modules powered by crank batteries (~2 tok/s), we introduce the *Monolith MN1 IPU*. Fabricated on sub-nanometer nodes with 3D CoWoS packaging and direct-to-chip micro-fluidic cooling, the MN1 eliminates the digital program counter entirely, enabling standalone, air-gapped physical appliances to continuously simulate multi-modal world states under a 400W thermal envelope.

Index Terms—World Models, Standalone Inference Appliances, Hardware-Model Duality, IBM Sub-1nm Lithography, Silicon Photonics, Compute-In-Memory (CIM), Probabilistic Bits (p-Bits).

I. INTRODUCTION & TAXONOMY OF IPUS

"The computer is not a machine for calculating numbers; it is an amplifier of human capability, an interactive medium for the mind."
— M. Mitchell Waldrop, *The Dream Machine* (2001) [1]

Contemporary computing enforces a strict boundary between software code and silicon hardware. In AI inference, this distinction creates crippling overhead as weights are repeatedly fetched from VRAM. To support real-time continuous World Models—which simulate physical visual, spatial, and environmental dynamics—we propose **The Hardware-Model Duality ("The Flipping")**: *The computer is no longer running the model; the computer IS the model, and the model IS the computer.*

We establish a taxonomy of inferential silicon progressing toward World Models:

1) Text LLM IPUs

Accelerating discrete text-token prediction (e.g., Groq LPU, Cerebras WSE-3, Etched Sohu).

2) Image Generation IPUs

Dedicated hardware for spatial latent diffusion steps.

3) Video Generation IPUs

High-throughput temporal-spatial latent streaming engines.

4) Unified World Model IPUs

Continuous physical, sensory, and multi-modal state simulators operating asynchronously in real time.

II. THROUGHPUT BENCHMARKS & METRICS

While LLM throughput is measured in discrete *tokens per second*, World Model IPUs require spatial-temporal metrics: *Latent Frames Per Second (Latent FPS \$ge 120\$)*, *Continuous Latent Sampling Rate (\$Hz/ ext{ms}\$)*, and *Voxel Transformation Bandwidth (\$ ext{VTX}/s\$)*.

Platform	Architecture	LLM Throughput	Target Substrate
Mac Studio / Off-Grid Node	Unified ARM / GGUF	~2 - 15 t/s	Crank/Battery Edge [2]
NVIDIA H100 Cluster	4nm Digital GPU	~200 t/s	PCIe / HBM Cluster
Groq LPU	14nm SRAM LPU	~800 t/s	Deterministic SRAM
Cerebras WSE-3	5nm Wafer SRAM	~2,500 t/s	Wafer-Scale Mesh [3]
Etched Sohu	4nm ASIC (8-chip)	~500,000 t/s	Hardcoded Attention [4]
Monolith MN1	Sub-1nm AIMC/Photonics	1,000,000 t/s	World Model Appliance

III. LITERATURE & PHYSICAL FOUNDATIONS

The concept of physically embodying neural equations in hardware builds upon decades of seminal literature:

A. Neuromorphic Computing (Mead, 1990)

Carver Mead proved subthreshold CMOS physics directly instantiates neural differential equations without digital code execution [5].

B. Memristive In-Memory Computing (Ielmini & Wong, 2018)

Ielmini and Wong demonstrated that non-volatile resistive switching crossbars execute matrix-vector products physically in memory using Kirchhoff's laws [6].

C. NeuRRAM (Wan et al., Nature 2022)

NeuRRAM proved full-scale resistive RRAM compute-in-memory delivers software-equivalent inference directly inside memory cells [7].

D. IBM NorthPole (Modha et al., Science 2023)

IBM NorthPole eliminated off-chip memory access by co-locating compute and memory in dense spatial physical tiles [8].

IV. LITHOGRAPHY, PHOTONICS & MATERIALS

A. IBM Sub-1nm (0.7nm Nanostack) Lithography

IBM's sub-1nm 3D vertical transistor nanostack technology enables packing ~100 billion transistors onto a fingernail footprint, offering 50% higher power, 70% lower energy, and 40% SRAM scaling [9].

B. Integrated Silicon Photonics

On-chip optical waveguides, micro-ring resonators, and optical TSVs move signals at light speed (299,792 km/s) for zero-heat multi-terabit/s Inter-Inference Core Communication (IICC) [10].

C. Material Science Innovations Required

Scaling standalone IPUs requires material breakthroughs: 2D Transition Metal Dichalcogenides (TMDs) for sub-nm channel transport, Lithium Niobate (LiNbO₃) electro-optic modulators for photonics, and micro-fluidic direct-to-chip 3D liquid cooling manifolds for >400 W TDP dissipation.

D. Thermodynamic Probabilistic Bits (p-Bits)

p-Bits utilize stochastic Magnetic Tunnel Junctions (s-MTJs) operating at low energy barriers to perform natural Boltzmann sampling [11]:

$$P_i(t) = \text{ext}\{\text{sgn}\} \text{Big}(\text{anh}(I_i(t)) - r(t) \text{Big}) \text{quad [Eq. 1]}$$

V. THE MONOLITH MN1 SUBSTRATE




Fig. 1. Monolith MN1 IPU depicting sub-1nm SRAM, photonics, and continuous-wave CIM/p-bit crossbars.

Quantizing 1T parameter multi-modal World Models into 1.58-bit ternary states (weights $\in \{-1, 0, 1\}$) enables an 850W desktop appliance to execute continuous real-time world simulations using wall-plug electricity.

VI. CONCLUSION

Fulfilling Waldrop & Licklider's vision, the Monolith IPU establishes the ultimate duality: *The computer is the model, and the model is the computer*. By combining sub-1nm nanostack lithography, silicon photonics, and in-memory computing, IPUs pave the way for autonomous standalone World Model appliances.

REFERENCES

1. M. M. Waldrop, *The Dream Machine*, Viking, 2001.
2. J. C. R. Licklider, "Man-Computer Symbiosis," *IRE HFE*, 1960.
3. Cerebras Systems, "WSE-3 Wafer Engine Architecture," 2024.
4. Etched AI, "Sohu: Transformer ASIC Infrastructure," 2024.
5. C. Mead, "Neuromorphic Electronic Systems," *Proc. IEEE*, 1990.
6. D. Ielmini & H. Wong, "In-memory computing," *Nat. Electron.*, 2018.
7. W. Wan et al., "NeuRRAM CIM chip," *Nature*, 2022.
8. D. Modha et al., "IBM NorthPole," *Science*, 2023.
9. IBM Research, "Sub-1nm Nanostack Chip Technology," 2026.
10. L. Chen et al., "Silicon Photonics for AI," *Nat. Photonics*, 2024.
11. C. Tyagi et al., "Probabilistic Computing," *Nat. Electron.*, 2024.
12. Y. LeCun, "A Path Towards Autonomous Machine Intelligence," 2022.
13. D. Ha & J. Schmidhuber, "World Models," *arXiv:1803.10122*, 2018.
14. Y. Hailu, "Capital Superintelligence Notes," 8T Capital, 2026.